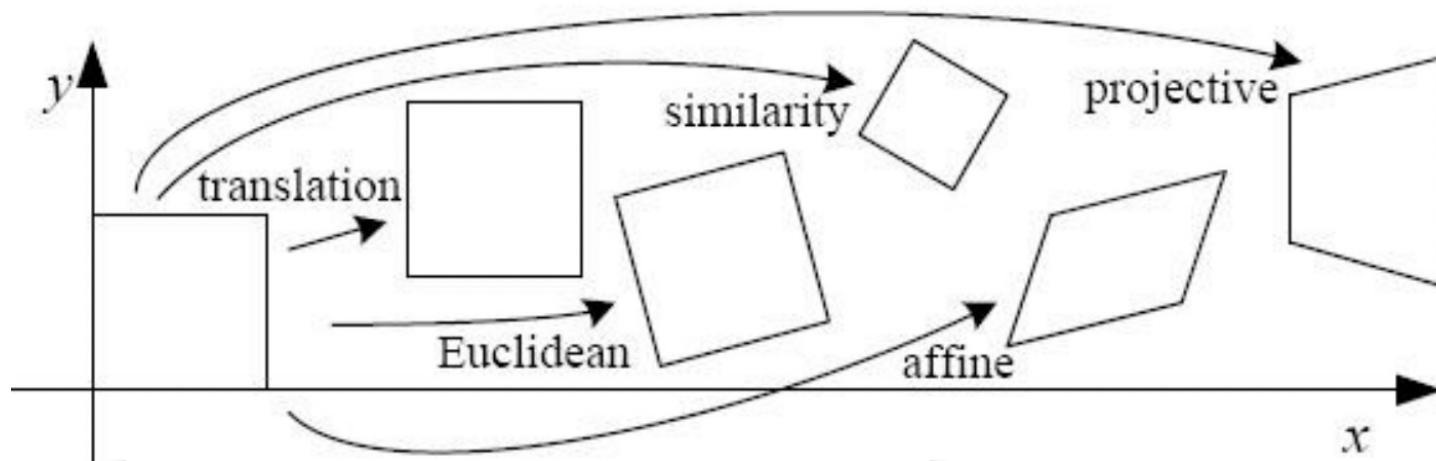


3. Image Transformations and Building Panaromas

Instructor: Shaifali Parashar

Many Slides from Fei-Fei Li's lectures at Standford

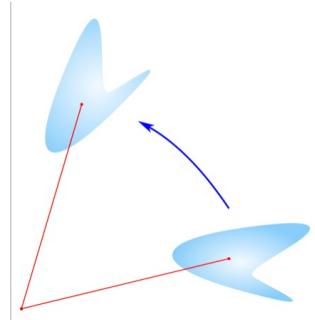
Transformations



Matrix Transformations: Rotation

In 2D: degrees of freedom = 1

$$\mathbf{R} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$



In 3D: degrees of freedom = 3

$$\mathbf{R}_x = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \phi & -\sin \phi \\ 0 & \sin \phi & \cos \phi \end{bmatrix}$$

$$\mathbf{R}_y = \begin{bmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{bmatrix}$$

$$\mathbf{R}_z = \begin{bmatrix} \cos \psi & -\sin \psi & 0 \\ \sin \psi & \cos \psi & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$\mathbf{R} = \mathbf{R}_z \mathbf{R}_y \mathbf{R}_x = \begin{bmatrix} \cos \theta \cos \psi & -\cos \phi \sin \psi + \sin \phi \sin \theta \cos \psi & \sin \phi \sin \psi + \cos \phi \sin \theta \cos \psi \\ \cos \theta \sin \psi & \cos \phi \cos \psi + \sin \phi \sin \theta \sin \psi & -\sin \phi \cos \psi + \cos \phi \sin \theta \sin \psi \\ -\sin \theta & \sin \phi \cos \theta & \cos \phi \cos \theta \end{bmatrix}$$

Rotation is an orthogonal transformation, $\det(\mathbf{R}) = 1$

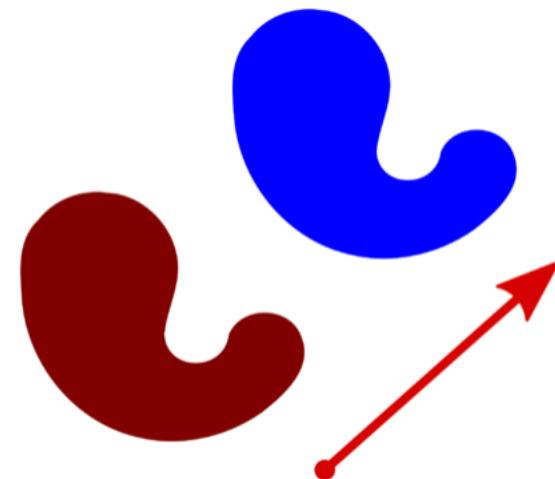
Reflection is like rotation, but $\det(\mathbf{R}) = -1$

Axis-angle representations: Euler angles, Quaternions, Rodrigue's formula

Matrix Transformations: Translation

In 2D: degrees of freedom = 2

In 3D: degrees of freedom = 3



Matrix Transformations: Scaling

In 2D:

$$\mathbf{S} = \begin{pmatrix} s_x & 0 \\ 0 & s_y \end{pmatrix}$$

Isotropic: $s_x = s_y$

Anisotropic: $s_x \neq s_y$

degrees of freedom (isotropic): 1

degrees of freedom (anisotropic): 2



Isotropic



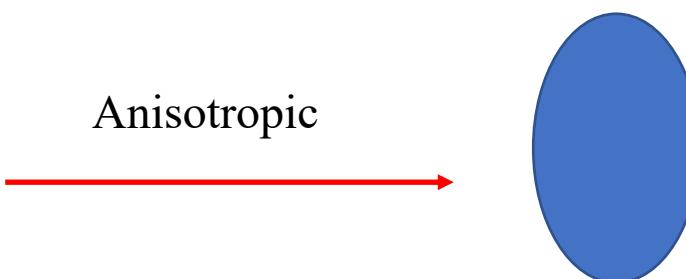
In 3D: $\mathbf{S} = \begin{pmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & s_z \end{pmatrix}$

$s_x = s_y = s_z$

$s_x \neq s_y \neq s_z$

degrees of freedom (isotropic): 1

degrees of freedom (anisotropic): 3



Matrix Transformations: Shear

In 2D:

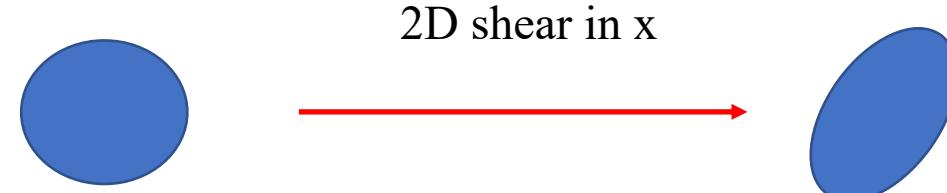
$$\mathbf{S} = \begin{pmatrix} 1 & s_x \\ s_y & 1 \end{pmatrix}$$

degrees of freedom: 2

In 3D:

$$\mathbf{S} = \begin{pmatrix} 1 & s_{x1} & s_{x2} \\ s_{y1} & 1 & s_{y2} \\ s_{z1} & s_{z2} & 1 \end{pmatrix}$$

degrees of freedom: 6



Euclidean Transformations

Also known as rigid transformations or isometries

Preserve lengths, angles and areas

A combination of rotations, reflections and translations

Euclidean Transformations

Also known as rigid transformations or isometries

Preserve lengths, angles and areas

A combination of rotations, reflections and translations

In 2D: degrees of freedom: 1(rotation) + 2(translation) = 3 $[R|t]_{2 \times 3}$

In 3D: degrees of freedom: 3(rotation) + 3(translation) = 6 $[R|t]_{3 \times 4}$

Similarity Transformations

Preserve ratio of lengths, angles

A combination of scale, rotations, reflections and translations

In 2D: degrees of freedom: $1(\text{scale}) + 1(\text{rotation}) + 2(\text{translation}) = 4$ $s[R|t]_{2 \times 3}$

In 3D: degrees of freedom: $1(\text{scale}) + 3(\text{rotation}) + 3(\text{translation}) = 6$ $s[R|t]_{3 \times 4}$

Affine Transformations

Preserve ratio of lengths, ratio of areas, parallelism

A combination of shear, scale, rotations, reflections and translations

In 2D: degrees of freedom: $2(\text{shear}) + 2(\text{scale}) + 1(\text{rotation}) + 2(\text{translation}) = 6$

In 3D: degrees of freedom: $6(\text{shear}) + 3(\text{scale}) + 3(\text{rotation}) + 3(\text{translation}) = 12$

Affine Transformations

Preserve ratio of lengths, ratio of areas, parallelism

A combination of shear, scale, rotations, reflections and translations

In 2D: degrees of freedom: $2(\text{shear}) + 2(\text{scale}) + 1(\text{rotation}) + 2(\text{translation}) = 6$

In 3D: degrees of freedom: $6(\text{shear}) + 3(\text{scale}) + 3(\text{rotation}) + 3(\text{translation}) = 12$

Degrees of freedom (dof) can't be more than matrix size.

In 2D, its $2 \times 3 = 6$ and in 3D, its $3 \times 4 = 12$.

Affine > Similarity > Euclidean

Transformations

To transform \mathbf{x} , we write (using homogeneous coordinates): $\bar{\mathbf{x}} = \mathbf{Ax} + \mathbf{t}$

$$\begin{pmatrix} \bar{\mathbf{x}} \\ 1 \end{pmatrix} = \begin{pmatrix} \mathbf{A} & \mathbf{t} \\ \mathbf{0}_{1 \times d}^\top & 1 \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ 1 \end{pmatrix}$$

d is dimension

Euclidean: $\mathbf{A} = \mathbf{R}$

Similarity: $\mathbf{A} = s\mathbf{R}$

Affine: $\mathbf{Ax} + \mathbf{t}$

Projective Transformations

Preserves collinearity

A combination of distortion, shear, scale, rotations, reflections and translations

In 2D: degrees of freedom: 8

In 3D: degrees of freedom: 15

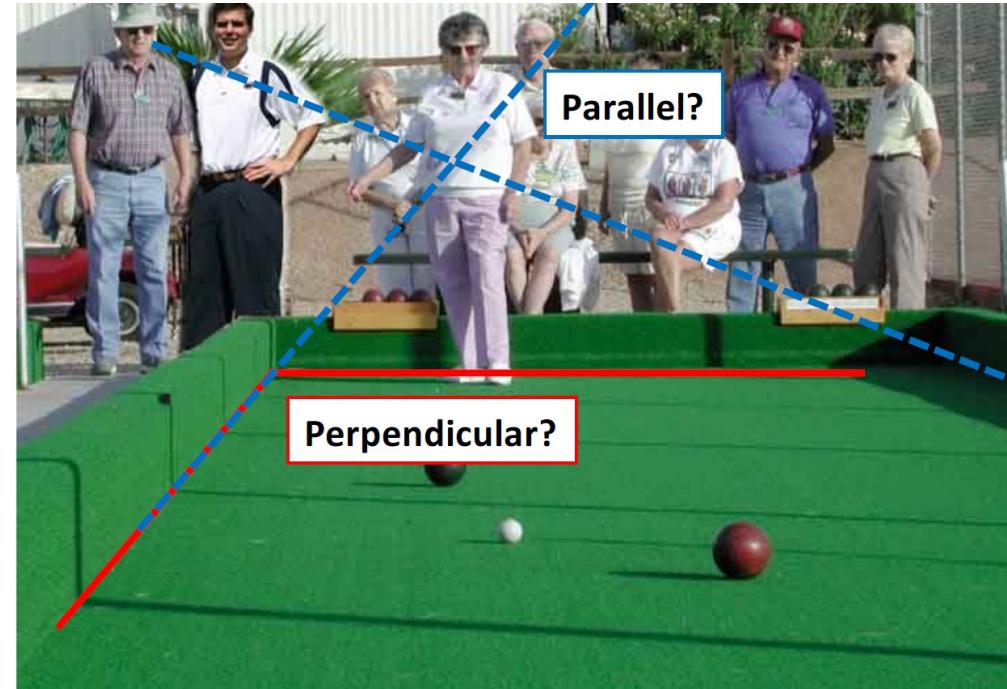
$$\begin{pmatrix} \bar{\mathbf{x}} \\ 1 \end{pmatrix} = \begin{pmatrix} \mathbf{A} & t \\ \mathbf{v}_{1 \times d}^\top & v \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ 1 \end{pmatrix}$$

Projective > Affine > Similarity > Euclidean

Always look at the rank of matrix to see how many dimensions does it impact

Images formation

Lengths, angles are lost
Parallelism is lost
Collinearity is preserved

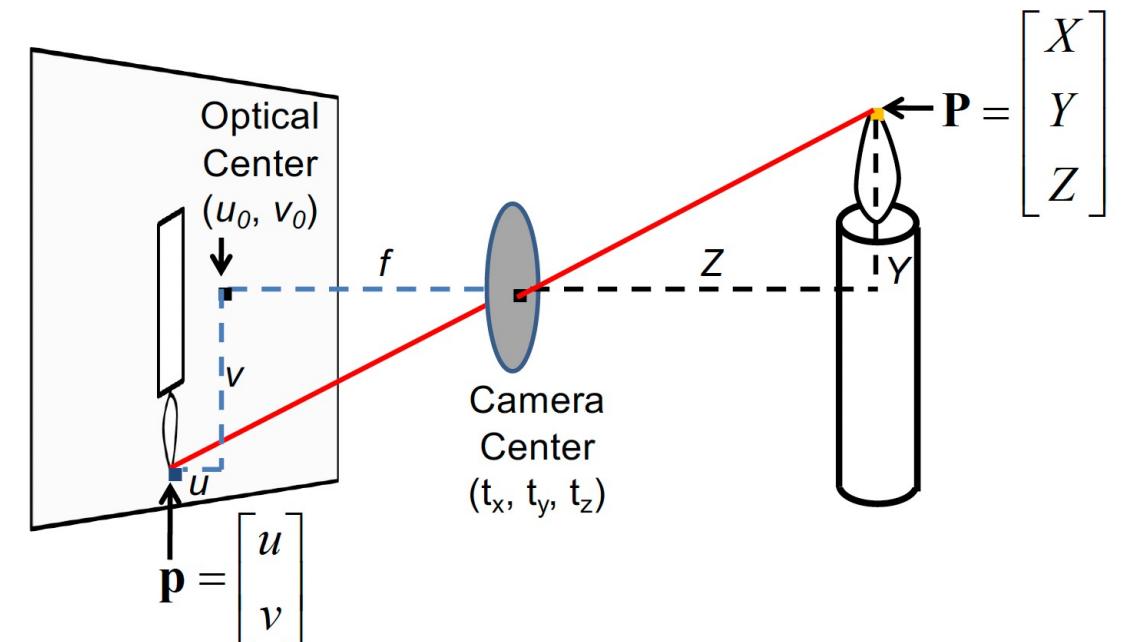


Images formation-Pin hole camera

$$\mathbf{p} = \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -f \frac{X}{Z} \\ -f \frac{Y}{Z} \end{bmatrix}$$

In fact, we invert and shift the image.

$$\mathbf{p} \longrightarrow [u_0, v_0] - \mathbf{p}$$



Homogeneous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

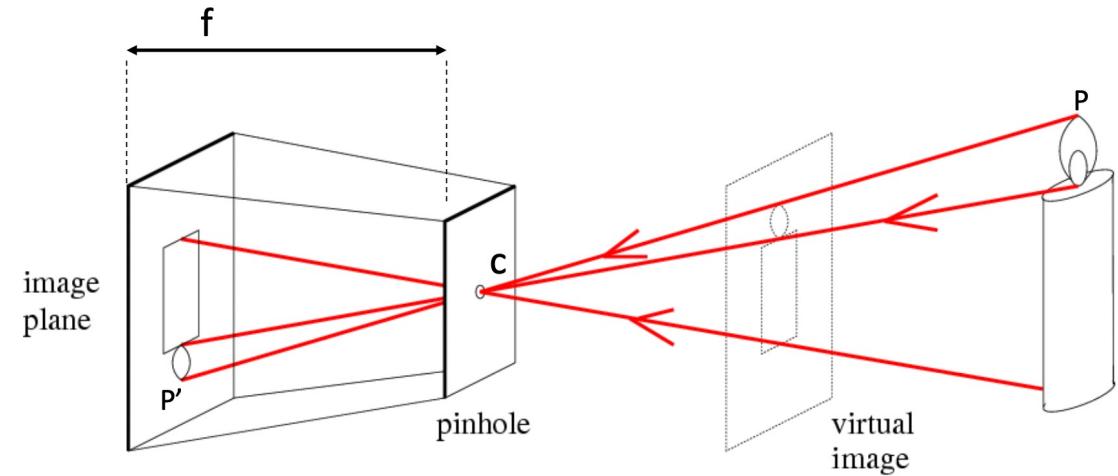
homogeneous scene
coordinates

Converting from homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

Also invariant to scale, multiplication with a scalar results movement along a ray



Another problem solved by homogeneous coordinates

Intersection of parallel lines

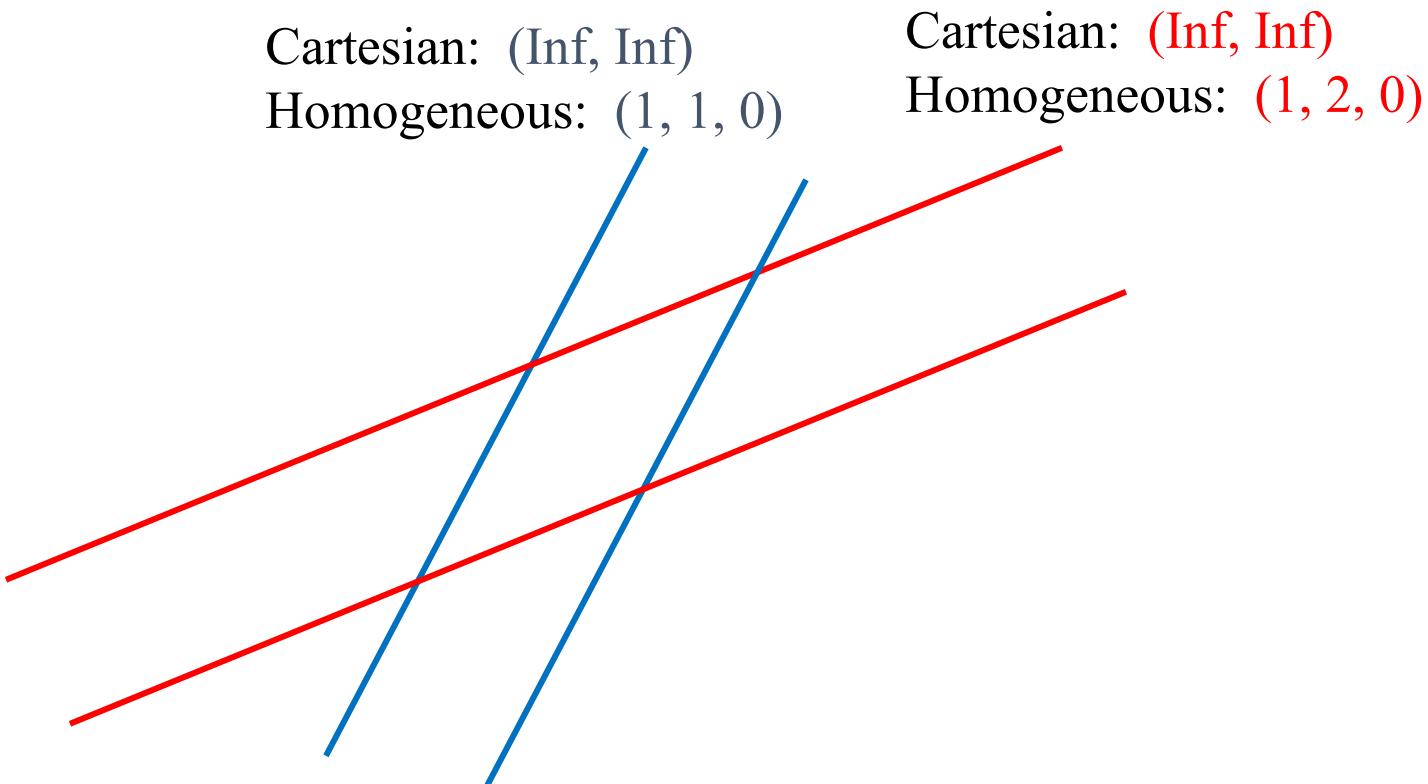
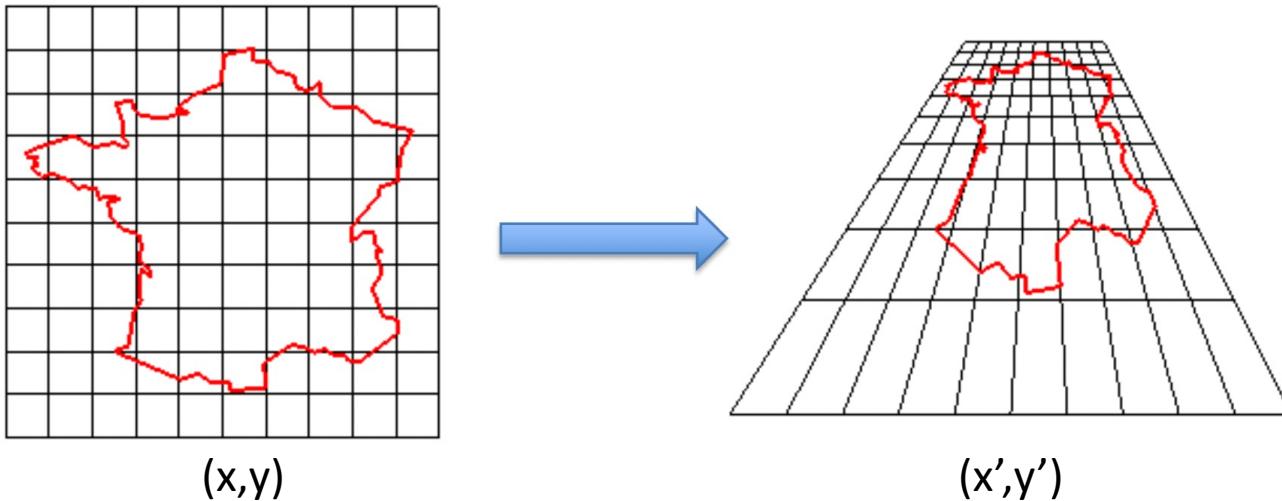


Image Transformations (Homography)



$$c \begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix}$$

Par I, Robert FERREOL, CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=2309125>

Image transformations → Projective transformation → Homography

Homography Estimation

- If point correspondences $(x,y) \leftrightarrow (x',y')$ are known

We can find a 3X3 transformation, \mathbf{A} such that

$$\mathbf{x}' = \mathbf{Ax}$$

$$\mathbf{A} = [a_1, a_2, a_3; a_4, a_5, a_6; a_7, a_8, a_9]$$

Remember from first lecture on transformations, a 2D projective transformation (most generic transformation) has 8 dof

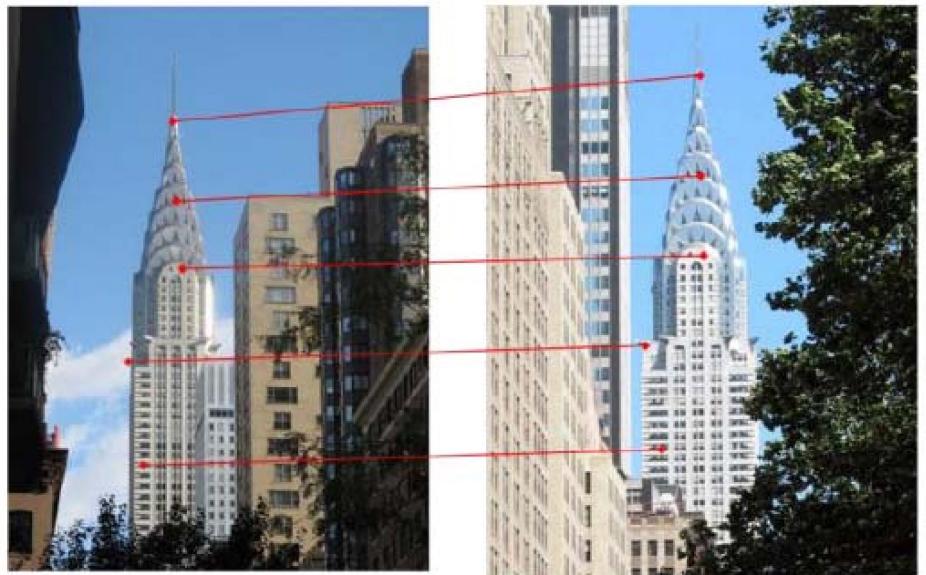
Assume $a_9 = 1$, therefore

$$x' = \frac{a_1x + a_2y + a_3}{a_7x + a_8y + 1}$$

$$y' = \frac{a_4x + a_5y + a_6}{a_7x + a_8y + 1}$$

$$\mathbf{x} = (x, y)$$

$$\mathbf{x}' = (x', y')$$



Homography Estimation

Assume $a_9 = 1$, therefore

$$x' = \frac{a_1x + a_2y + a_3}{a_7x + a_8y + 1}$$

$$y' = \frac{a_4x + a_5y + a_6}{a_7x + a_8y + 1}$$

$$a_7x'x + a_8x'y + x' = a_1x + a_2y + a_3$$

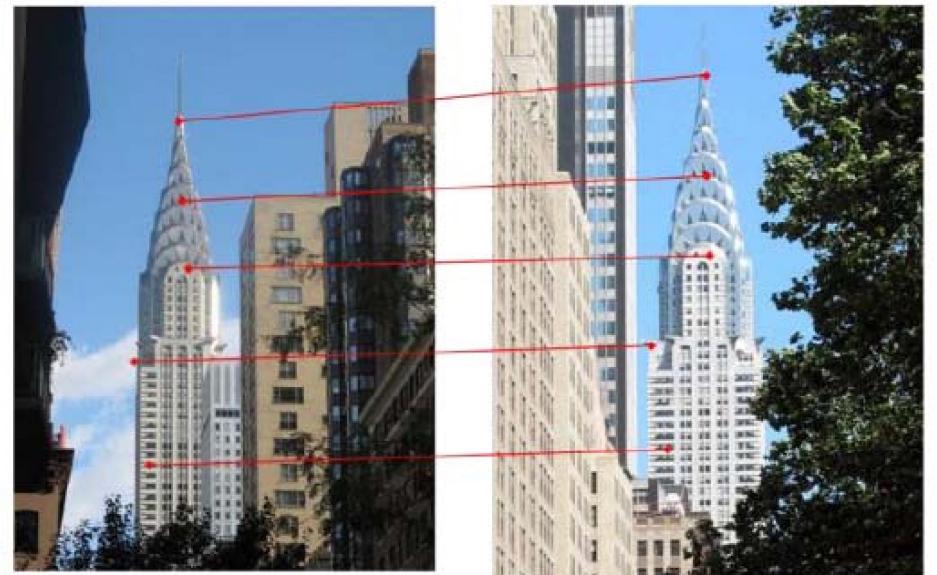
$$a_7y'x + a_8y'y + y' = a_7x + a_8y + a_6$$

$$x' = a_1x + a_2y + a_3 - a_7x'x - a_8x'y$$

$$y' = a_7x + a_8y + a_6 - a_7y'x - a_8y'y$$

$$\mathbf{x} = (x, y)$$

$$\mathbf{x}' = (x', y')$$



Segregate the equations into matrices/vectors of knowns (\mathbf{x} , \mathbf{x}') and unknowns (a_1-a_8)

Homography Estimation

$$a_7x'x + a_8x'y + x' = a_1x + a_2y + a_3$$

$$a_7y'x + a_8y'y + y' = a_7x + a_8y + a_6$$

$$x' = a_1x + a_2y + a_3 - a_7x'x - a_8x'y$$

$$y' = a_7x + a_8y + a_6 - a_7y'x - a_8y'y$$

Two rows for each point i

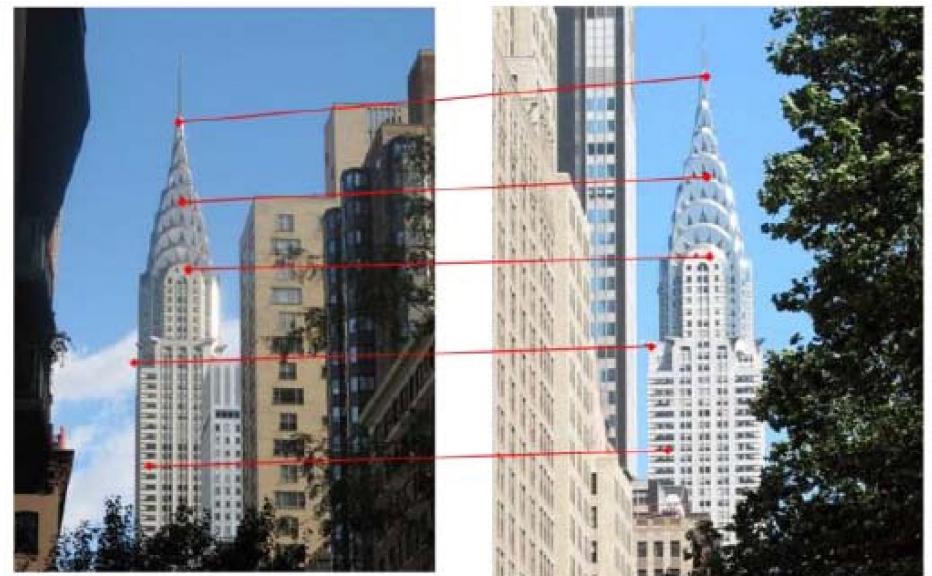
$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & \vdots & 0 & -x_i x'_i & -y_i x'_i \\ 0 & 0 & 0 & x_i & y_i & \vdots & 1 & -x_i y'_i & -y_i y'_i \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \\ a_7 \\ a_8 \end{bmatrix} = \begin{bmatrix} \vdots \\ x'_i \\ y'_i \\ \vdots \end{bmatrix}$$

$$\mathbf{T}_{8 \times 8} \mathbf{a}_{8 \times 1} = \mathbf{B}_{8 \times 1}$$

Linear Least squares: $\mathbf{a} = (\mathbf{T}^T \mathbf{T})^{-1} \mathbf{T}^T \mathbf{B}$

$$\mathbf{x} = (x, y)$$

$$\mathbf{x}' = (x', y')$$



Homography using Direct Linear Transformation (DLT)

Don't assume $a_9 = 1$, add it to the system

In order to obtain a non-trivial solution,
we obtain \mathbf{a} :

$\min \| \mathbf{T}'\mathbf{a} \|$ such that $\|\mathbf{a}\| = 1$

Two rows for each point i

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & \vdots & 0 & -x_i x'_i & -y_i x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -x_i y'_i & -y_i y'_i \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \\ a_7 \\ a_8 \\ a_9 \end{bmatrix} = \begin{bmatrix} \vdots \\ x'_i \\ y'_i \\ \vdots \end{bmatrix}$$

$$\mathbf{T}'_{8 \times 9} \mathbf{a}_{9 \times 1} = \mathbf{0}_{9 \times 1}$$

$$\mathbf{T}' = [\mathbf{T}, -\mathbf{B}]$$

Homography using Direct Linear Transformation (DLT)

Don't assume $a_9 = 1$, add it to the system

In order to obtain a non-trivial solution,
we obtain \mathbf{a} :

$$\min \| \mathbf{T}'\mathbf{a} \| \text{ such that } \|\mathbf{a}\| = 1$$

Solution: SVD of \mathbf{T}' . The eigenvector corresponding to smallest eigenvalue of $\mathbf{T}'^T\mathbf{T}'$ is the solution.

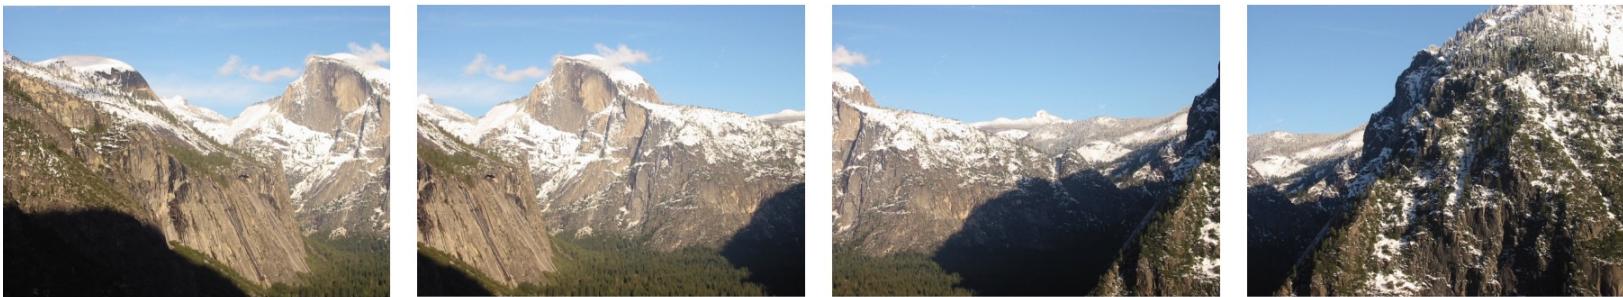
In order to get a stable solution, Normalise the data before DLT.

- a) Translate for zero mean
- b) Scale so that average distance to origin is $\sim \sqrt{2}$

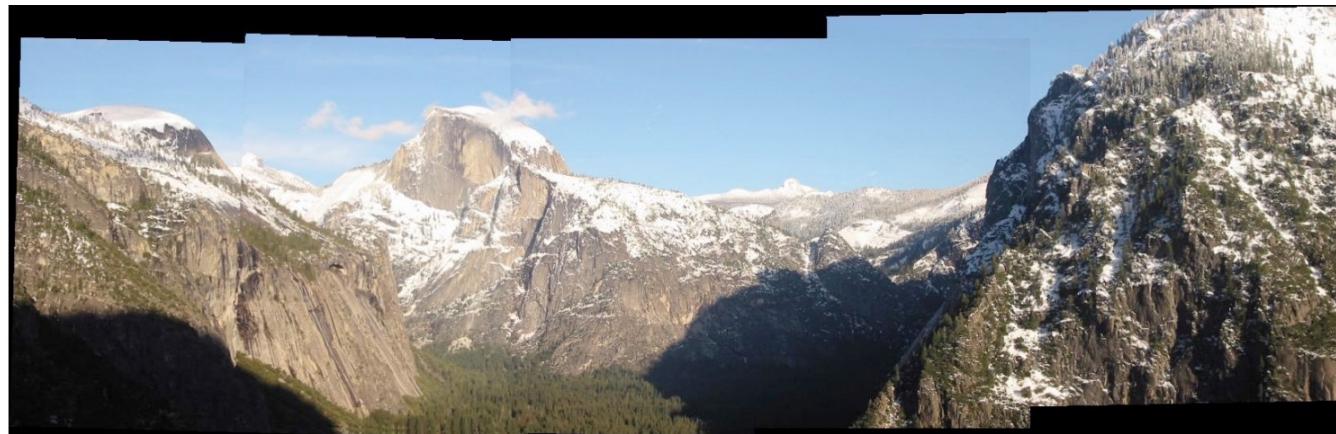
https://www.cs.cmu.edu/~16385/s17/Slides/10.2_2D_Alignment_DLT.pdf

Panoramic Image Stitching

Given a set of images



Obtain:



Problem with Linear Least Squares (LLS) fitting

1. Noisy data and outliers

Solution:

- a) Voting. Check all possible combinations. Cycle through features, cast votes for model parameters. Bad idea due to high computation cost.
- b) RANdom SAmple Consensus (RANSAC)

RANSAC

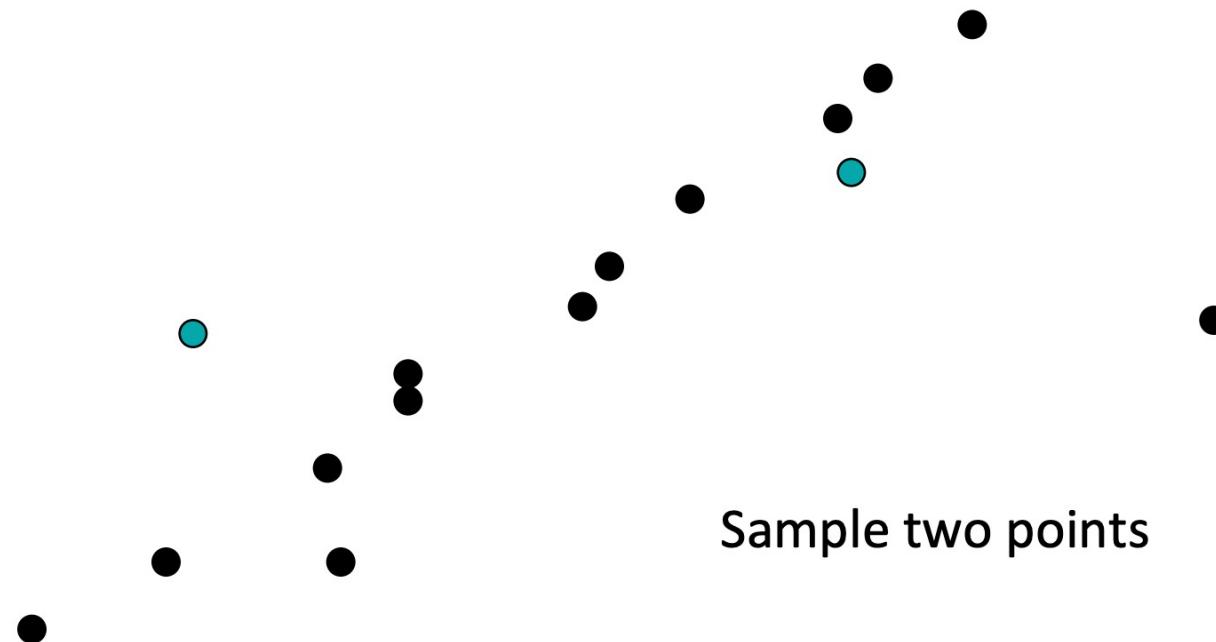
Look for ‘only’ inliers.

1. Randomly select a seed group to estimate transformation
2. Compute transformation
3. Compute inliers
4. If inliers are large, recompute transformation

Keep transformation with most inliers

RANSAC: line fitting example

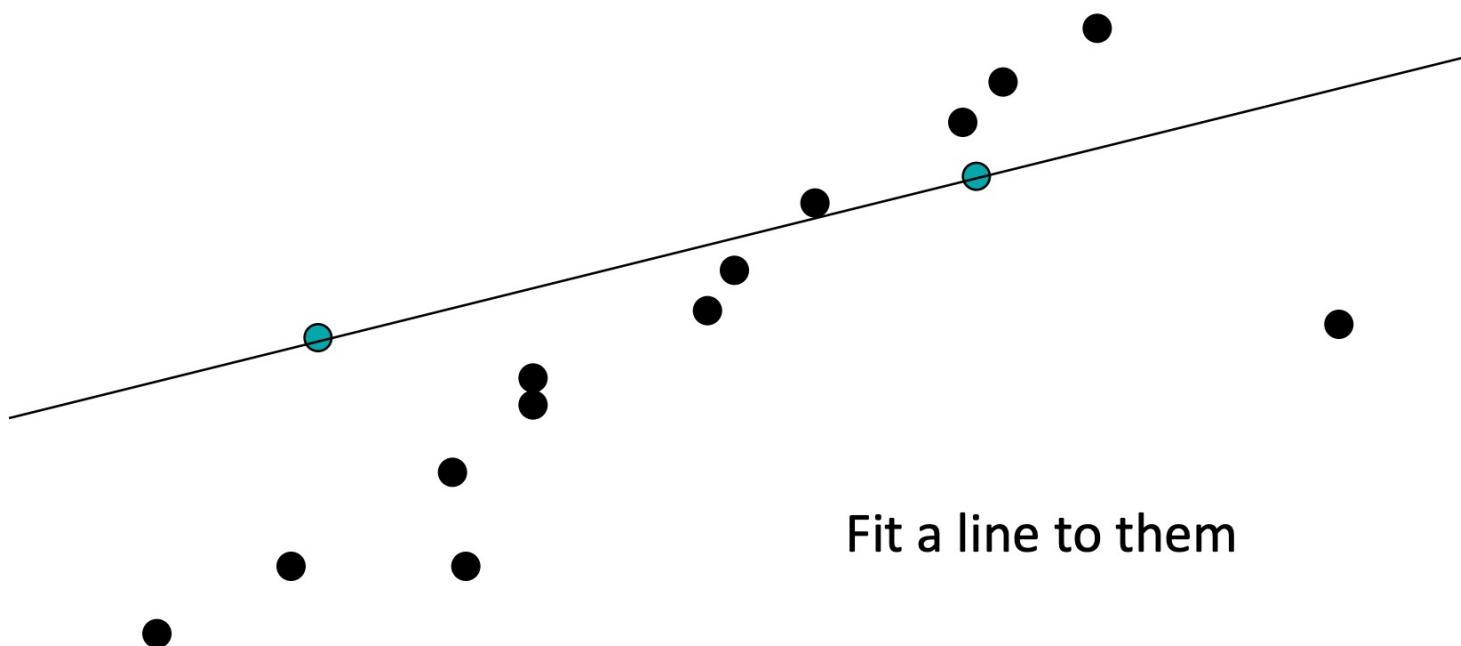
Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line fitting example

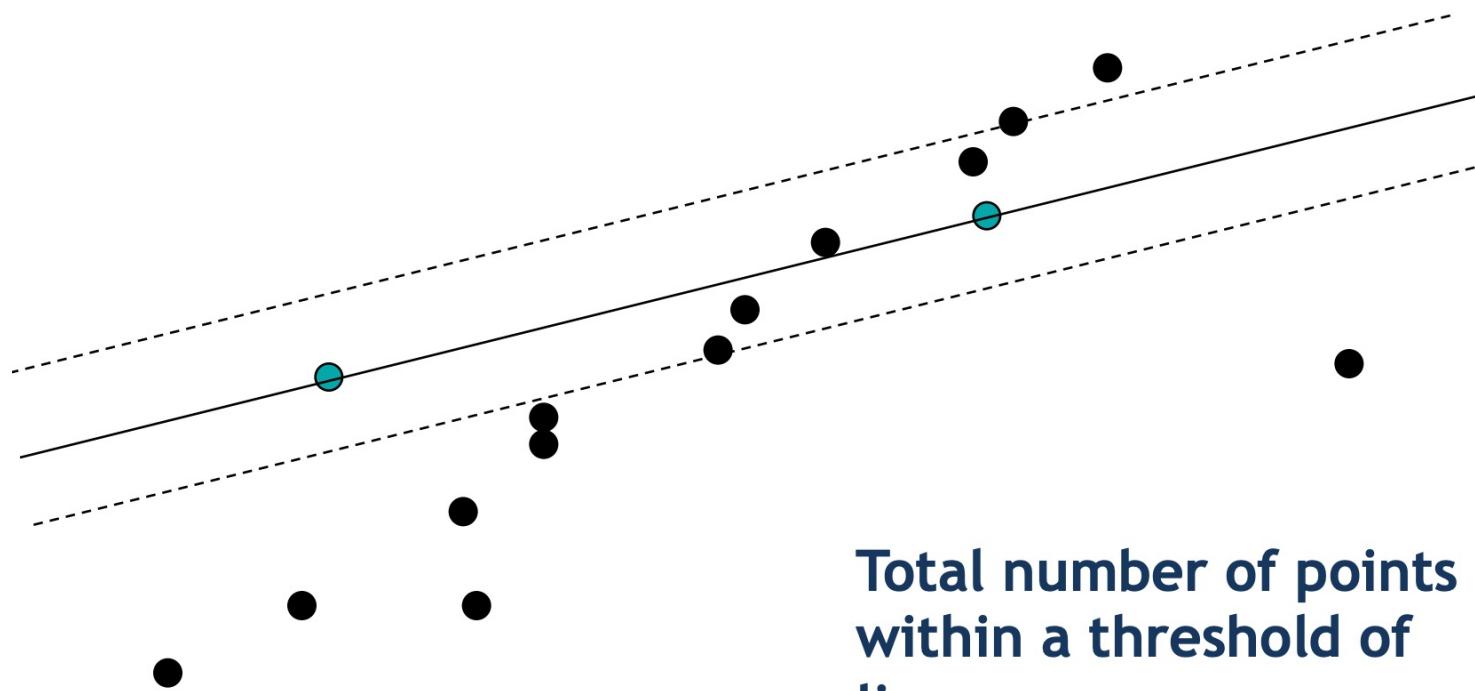
Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line fitting example

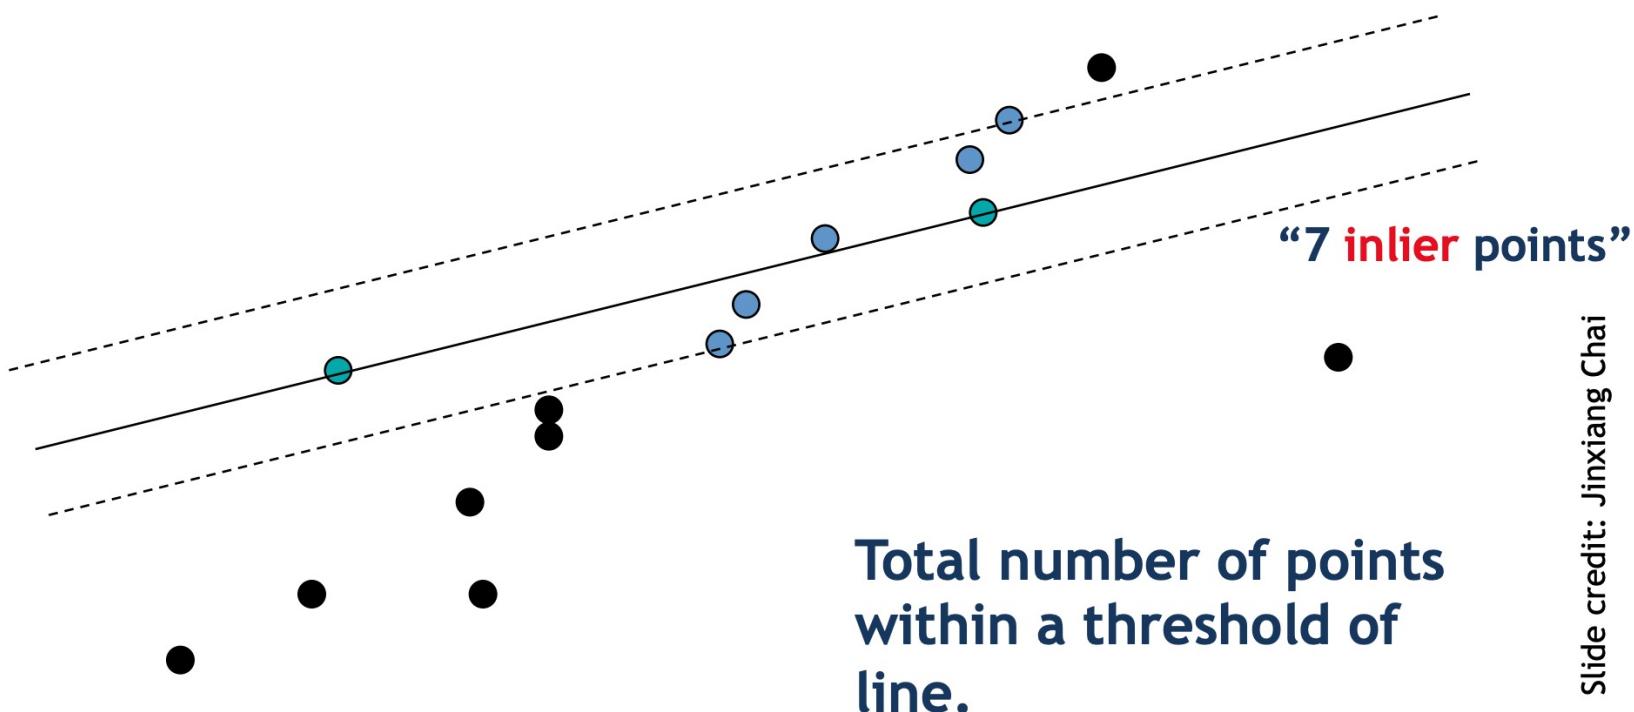
Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line fitting example

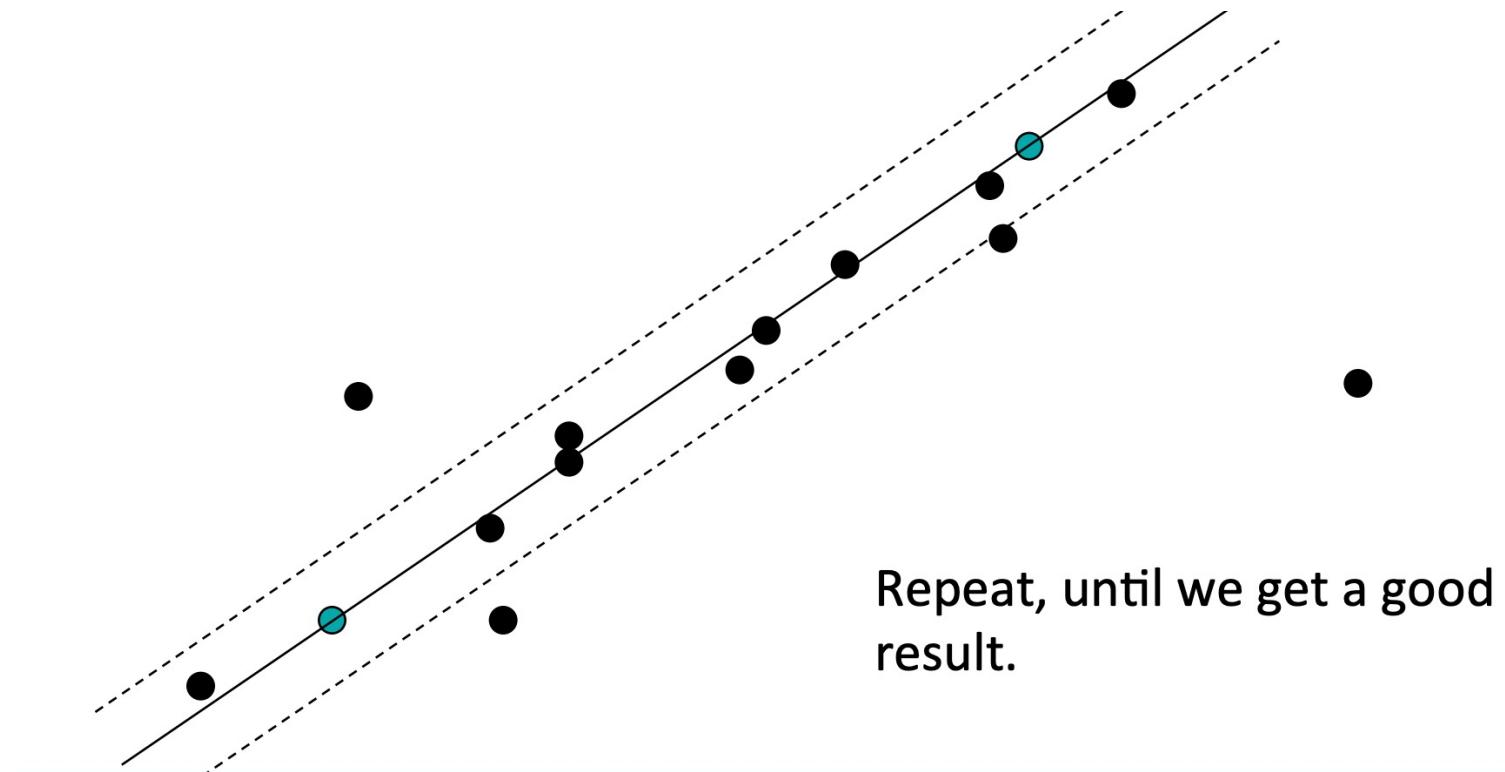
Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line fitting example

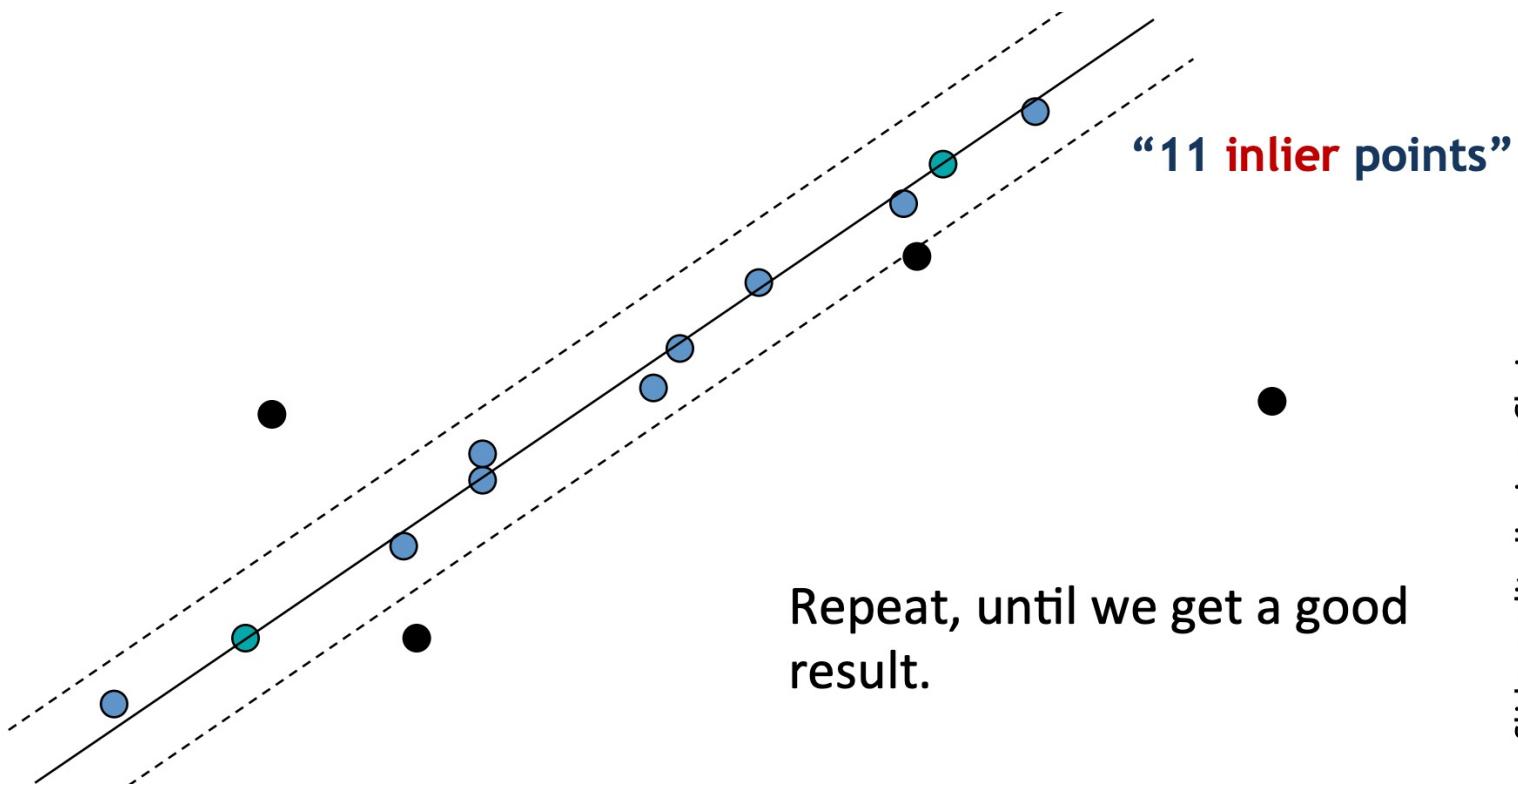
Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line fitting example

Estimate the best line



Slide credit: Jinxiang Chai

RANSAC: line-fitting example

After RANSAC, use all inliers to improve transformation. This may change the transformation, so perform another reclassification of inliers-outliers.

RANSAC is fast but it can accommodate relatively small number of outliers.

Voting is efficient but computationally expensive.

RANSAC: line-fitting example

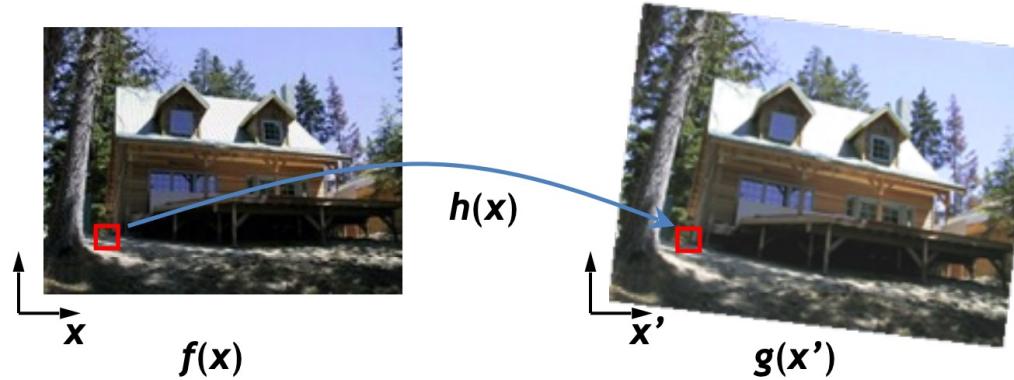
After RANSAC, use all inliers to improve transformation. This may change the transformation, so perform another reclassification of inliers-outliers.

RANSAC is fast but it can accommodate relatively small number of outliers.

Voting is efficient but computationally expensive.

Image Warping: Forward

Given a transformation $x' = h(x)$ and source image $f(x)$, how do we compute a transformed image $g(x') = f(h(x))$?



- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels, normalize later (*splatting*)

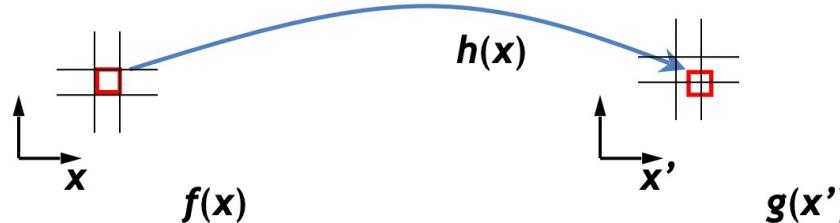
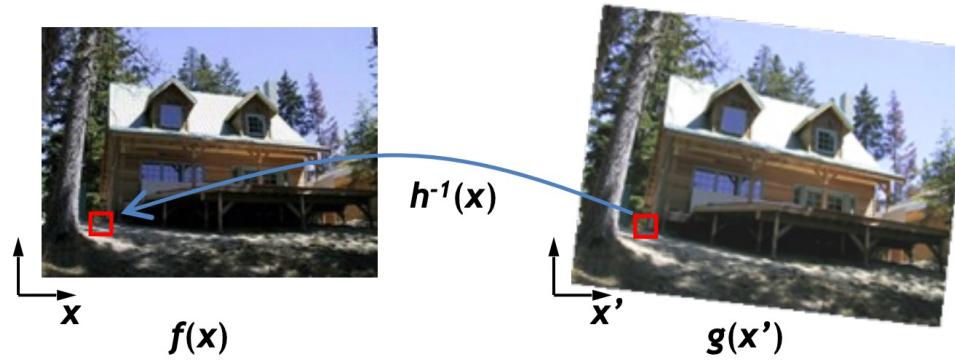
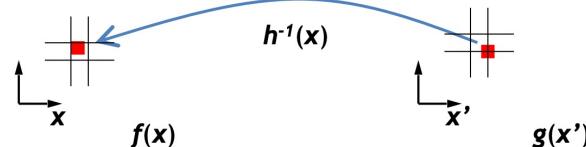


Image Warping: Backward

Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$?

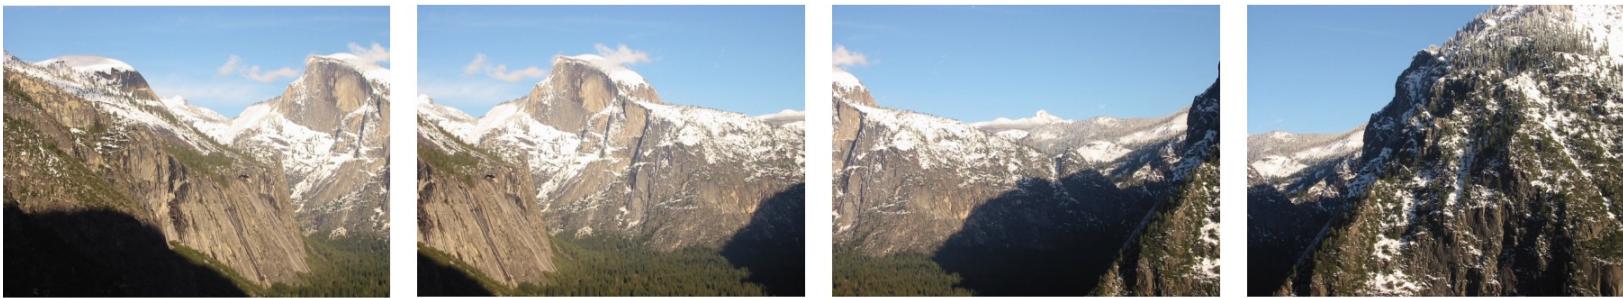


- What if pixel comes from “between” two pixels?
- Answer: *resample* color value from *interpolated* source image



Panoramic Image Stitching

Given a set of images



Obtain:

